

I-EPOS Manual

Peter Pilgerstorfer

December 17, 2016

1 Introduction

1.1 Install and setup

TODO: netbeans/eclipse instructions TODO: add libraries

1.2 Execute the sample simulation

TODO: run GUI TODO: run from code: SimpleExperiment.main

2 Architecture

TODO: software architecture

3 Use cases

3.1 Configure the simulation

Let's take a look at the provided sample experiment SimpleExperiment. It contains a main function that specifies all parameters, starts the simulation and presents the results.

Dataset

Initially the dataset is specified. Note that the Dataset object is not required for the algorithm, but it simplifies the configuration. Technically all we need is a way to get a list of possible plans for each agent. The Dataset interface provides this functionality via Dataset.getPlans(int agentIdx). There are two classes of datasets implemented:

- FileVectorDataset is a dataset that is read from disk. Only the dataset folder has to be specified. The example datasets are located in the directory 'project dir./input-data'. Section 3.4 describes the input format for this kind of dataset.

- GaussianDataset is a generated dataset where every plan is a vector drawn from a gaussian distribution. The parameters specify the number and dimensionality of the plans as well as mean and standard deviation of the distribution.

The number of agents can be set arbitrarily. However, when using a FileVectorDataset, the number of agents has to be below FileVectorDataset.getNumAgents().

Cost functions

The global cost function describes what we want to minimize globally. The local cost function describes what each agent wants to minimize locally. λ is the tradeoff between global and local minimization. $\lambda = 0$ means only global cost is minimized, $\lambda = 1$ means only local cost is minimized. For global cost functions we can choose any implementation of the interface CostFunction in general. However, for gradient descent based algorithms an instance of DifferentiableCostFunction is required. A list of possible cost functions is as follows:

- DotCostFunction minimizes the dot product of a given vector with the global response. See Section .2.1 how to read a vector from a file. Be aware that the dimensionality of the vector has to match the dimensionality of the dataset. An example use case for this cost function is minimizing monetary cost. The agent plans contain the amount of resources they consume, and the vector passed to DotCostFunction describes the price for each resource. As this is a linear cost function, I-EPOS always finds the optimal value in the first iteration.
- SqrDistCostFunction minimizes the (squared) distance of the global response to a given vector. See Section .2.1 how to read a vector from a file. Be aware that the dimensionality of the vector has to match the dimensionality of the dataset. This cost function tries to make the global response as similar to the provided target vector as possible.
- VarCostFunction minimizes the variance of the global response. Therefore it can be used to stabilize resource consumption over time, or for load balancing applications.
- StdDevCostFunction minimizes the standard deviation of the global response. For I-EPOS there is no difference between minimizing standard deviation and minimizing variance, as the functions share the same minima.
- MaxCostFunction (non-differentiable) minimizes the maximum value of the global response. This function is useful for peak reduction.

Local cost functions have to implement the PlanCostFunction interface. Two functions are implemented:

- IndexCostFunction lets agents select plans with a small index. Therefore the plans in a dataset should be ordered in a way that lists preferable plans first.

- `PlanScoreCostFunction` lets agents select plans with a small score. The dataset specifies the score for each plan. See Section 3.4 for details how to specify this information in a dataset.

Network

The network is considered to be a balanced tree with the same number of children for each inner node. The number of children for inner nodes can be specified. Be aware that the runtime of I-EPOS is exponential in the number of children. Luckily the optimization performance of a binary tree is already close to optimal in practice.

Logging

Next we specify what information we want to gather from the simulation. For this task we specify a `LoggingProvider[A]`, where `A` is the class of the agent that is used in the simulation. We then specify all information we want to log by adding `AgentLoggingProviders`. Each `AgentLoggingProvider` is responsible for reading and presenting one specific type of data. The output is presented after the simulation by calling the method `LoggingProvider.print`. The following loggers are ready to use:

- `GlobalCostLogger` logs the global cost in each iteration and prints the global cost for each individual iteration averaged over multiple simulations. Note that the sample shown in `SimpleExperiment` only performs one simulation. Multiple simulations can be performed with e.g. different seeds for the agents or different datasets. The only requirement is that the same `LoggingProvider` is used.
- `LocalCostLogger` logs the average local cost in each iteration and prints the average local cost for each individual iteration averaged over multiple simulations.
- `TerminationLogger` logs how many iterations it took for the algorithm to terminate. The algorithm is considered terminated if nothing changes between two consecutive iterations.
- `ProgressLogger` prints symbols every few iterations in order to show how far the algorithm has proceeded. It is only useful for large simulations.
- `JFreeChartLogger` shows a plot with global and (optionally) local cost values for each iteration in a new window. The logger requires `GlobalCostLogger` to be added to the `LoggingProvider` as well. If the `LocalCostLogger` is present, the local cost is also shown in the plot.
- `GraphLogger` shows a graph of the network at a given iteration in a new window. With the arrow keys you can switch between different iterations. Each agent is represented in a certain color. The color code depends on the specified type `GraphLogger.Type`. The following types are available:

- Change marks each agent that changed its selection in the previous iteration as black and all agents without change as white.
- Index colors each agent based on the index of the selected plan. Agents that selected the plan with minimal index are colored white and agents that selected the plan with maximal index are colored black.
- FileWriter writes the log to the specified directory once LoggingProvider.print() is executed.
- FileReader reads the log from the specified directory once LoggingProvider.print() is executed. This logger can be used to show results from a previous simulation that were stored with FileWriter. A sample application of FileReader can be seen in ReplayExperiment.

Algorithm

Finally, we have to specify the optimization algorithm. The algorithm is determined by the type of Agent that is used.

- IeposAgent has quite a few options that were part of the research. We need to specify the number of iterations the algorithm should perform. For problems with less than 1000 agents the (local) optimum is usually found with less than 20 iterations. In addition we can specify a PlanSelector. The default is IeposPlanSelector. As part of the research that was done for I-EPOS, gradient descent motivated plan selectors were also developed. However, they are inferior to the default in terms of optimization performance.
- CohdaAgent is an algorithm that was used as a baseline for I-EPOS. We only need to specify the number of iterations for this algorithm. Limitations: First, CohdaAgent does not support local cost. Therefore LocalCostLogger cannot be used. Second, the algorithm starts with an incomplete global response that is missing data from some agents. It takes $\log(\text{numAgents})/\log(\text{numChildren})$ iterations for the global response to be complete. Third, even though COHDA does not require the network to be a tree, only tree networks can be simulated with this software.

3.2 How to store evaluation results

Evaluation results can be stored using the LoggingProvider FileWriter. Once the print command is executed on the LoggingProvider instance, the log is written to the specified file. With FileReader, the written file can be read again. See ReplayExperiment as an example how to read a log file.

3.3 Write a new cost function

Write a new class that extends the abstract class `CostFunction<DT>` or `DifferentiableCostFunction<DT>` where DT is the datatype that this cost function should operate on.

The function `CostFunction.calcCost(DT value)` should compute the cost of the given value. For differentiable functions we also need to implement `DifferentiableCostFunction.calcGradient(DT value)` that should return the gradient of the function at point value.

3.4 Add a new dataset

One way of adding a new dataset is to use the existing class `FileVectorDataset` to read a custom dataset from the dataset directory. The dataset is a directory containing one file for each agent. The files should be named `agent_{id}.plans`, where `{id}` is the id of the agent, starting from 0 upwards. Each file should be a text file containing one row for each possible plan the agent can choose. A plan has the following layout: `{score};{vector}`. The score is a double value that describes the cost this plan imposes for an agent. It can be used for local cost minimization¹. The vector is a comma separated list of double values.

It is also possible to code a new dataset. The only requirement for the dataset is to generate a list of plans given the index of an agent. `Dataset` is a handy interface that can be used to implement a new dataset. While the sample datasets all use vectors as datatype, the new dataset can use a custom datatype.

3.5 Add a new datatype

Write a new class that implements the interface `Datatype`. The functions of the interface `Datatype` were designed work for vectors. So when implementing those functions, be aware that the semantic should be as it would be for vectors.

To use the new datatype, we also need to implement a new dataset and a new cost function that can handle the new datatype.

.1 Glossary

.2 Utility functions

.2.1 Reading a vector

A vector can be read from a file via `VectorIO.readVector(File vectorFile)`. The file is assumed to be text file, containing a comma-separated list of double values that make up the vector.

¹Set `lambda=0` if the score should be ignored.